

# Matrix Bubbling

Adam Busis, Jacky Lee, Princewill Okoroafor

November 2, 2018

## 1 Abstract

Matrix bubbling is a useful technique for analyzing big data sets because it collapses the large amount of data given into a smaller matrix of summary statistics, which is easier to work with. In this paper, we discuss matrix bubbling in two dimensions but this technique can be easier extended to higher-dimensional space.

## 2 Background

When working with big data, it often takes a long time to process data since there is so much of it. To try to resolve this problem, we want to be able to somehow reduce this data into its essence such that we numerically have less data to work with, which means we can process it faster, but still have sufficient information to extract something meaningful out of it.

One method that does this is called dimension reduction. The general idea of dimension reduction is that we can reduce the number of variables involved in our data such that we have less to work with but still enough information to perform computations with. Dimension reduction is a useful tool that enables us to analyze big data sets with less computation time.

Generally, this is done by determining which features define the data the most and restricting our attention to those few features. For example, consider the scenario where we want to read characters off of hand-drawn maps. Given a random symbol from the map, how do we determine whether the symbol is a letter or just some junk from the map? There can be a lot of information hidden within each symbol. One thing we could do is to only consider specific features that seem to be the most important for determining whether a symbol is a letter or not. For example, we can only look at the width, height, and density of the input and from there guess whether the symbol is a letter or not. This is

plausible because letters on a map generally are within some size limit; they are probably not too big (taking up a quarter of the map) and they are probably not too small (taking up just 3 pixels worth of space).

Another thing that can be done is to define new features that can be thought of as combinations of multiple features. For example, we could replace our width and height feature for a new perimeter feature, which is computed by taking the sum of the width and height then multiplying by 2. We see that this perimeter feature is a combination of both the width and height features.

Currently, the main dimension reduction technique is called Principle Component Analysis, or PCA for short. Essentially, this method takes in data and computes the most important features that help identify the data.

### 3 Motivation

Let us restrict our attention to the two-dimensional case for now since it is easier to visualize. Suppose we have a large set of data distributed in two-dimensional space in a fashion that is not Gaussian. For example, the pixels representing the letter ‘A’ are not distributed normally. Now suppose we would like to compare it with another data set that is also large and distributed in some space. Let’s call these two data sets  $A$  and  $B$ . One way to do it is to take each data point’s location in the space and compare it to the analogous space in the other data set. This requires going through each point and requires a lot of computations.

The idea behind matrix bubbling is that we can somehow summarize each data set using some proxy such that if we want to compare two data sets  $A$  and  $B$ , we can simply compare their proxies,  $A_p$  and  $B_p$ , where the proxies are much smaller than the original data sets. The “proxy” can be any smaller set of data depending on the data set. This can be thought of as a sort of dimension reduction since we are reducing the amount of data we need to consider and instead summarizing the data using a few key statistics. Ideally, the proxies for two different data sets will be close together if (and only if) the two data sets are similar, where similarity of the data sets is a criterion that depends on the application.

One possible way of summarizing the data would be to look at summary statistics for the entire data set, such as the mean and variance. The problem with this method is that just the mean and variance for the entire data set don’t have enough information to capture the nuances of the data set. The key idea is that we can view our space as a grid and partition this grid into smaller grids. Then for each grid, we can find a probability distribution that summarizes the

data within that grid and use it as a proxy. Since probability distributions can be summarized by a few statistics such as the mean and covariance for a Gaussian distribution, this reduces the number of data points we have to consider. By changing the number of smaller grids (and therefore the size of the smaller grids) that we divide the data set into, we can adjust the fineness of the details we look at in the data based on the application.

## 4 Introduction

Suppose we have a matrix  $M$  representing a two-dimensional data set which we would like to compare to another matrix  $N$  representing a different data set. To do this comparison, we would want to use a metric. Some of the available options for this include euclidean distance, Lp norms, etc. Any of these will work fine if our data set was small. But what happens when the data sets get very large? Pairwise distance metrics get very difficult to compute, and biases become able to skew the distances off significantly. Because of how applicable big data is, these are often the scenarios we run into the real world.

To solve this problem of comparing two matrices  $M$  and  $N$  representing two dimensional data sets, we propose matrix bubbling.

The first step to matrix bubbling involves summary. The goal of summary is to make the large data set finite and computationally small while retaining all the essential features of the data set. Summary sets out to transform the original  $m$  by  $m$  matrix where  $m$  is significantly large to an  $n$  by  $n$  matrix where  $n$  is computationally small. As this point, one might start to wonder - how large is significantly large and how small is computationally small - and the answer to this is - it depends. It depends on the data set and the application for which the data sets are used. In our tests of this approach, we transform high quality images of about  $1920 \times 1024$  to low quality images of about  $288 \times 288$ . Once the dimensions have been chosen, the goal would be to divide the original  $m$  by  $m$  matrix to square grids of size about  $\frac{m}{n}$ . A pictorial representation of this has been put below.

Then for each square grid, we fit all the points in the square grid to a normal distribution. This reduces all these points to two values( $\mu, \sigma$ ) that summarize them.

Once we've summarized both matrices  $M$  and  $N$  to  $M'$  and  $N'$  with lower dimensions, we can now proceed to compare. As a result of the nature of our reduction (with probability distributions), we cannot simply use any of the regular distance metrics like euclidean, Lp norm, etc. This is because we could have two normal distributions that are the same but orders of magnitude away




Figure 1: Visual Representation of a Bubbled Grid[1]

from each other by euclidean distance, etc so we need a theoretical sound way of comparing probability distributions.

## 5 Methods

For our data set, we started with a set of black cursive letters.



ALPHABETPRINTABLES.ORG

Figure 2: Sample Cursive Letter

We then tried to apply some morphological operators to the letter ‘a’ to perturb the image. This was so that we could compare this image with the original image and see whether the scores indicated that the images were similar. We also compared them to different letters to see if the scores indicated the images were different.

### 5.0.1 Morphological Operators

Morphological operations apply a structuring element to an input image, creating an output image of the same size. In a morphological operation, the value of each pixel in the output image is based on a comparison of the corresponding pixel in the input image with its neighbors. By choosing the size and shape of the neighborhood, you can construct a morphological operation that is sensitive to specific shapes in the input image.[2]

We'll use the following morphological operations to try to mimic real life transformations to images. A lot of work is currently being done in the area of interpreting historical maps. Historical maps help us understand how geographical locations change. Because of how old some historical maps are, they may have undergone various transformations. Dilation and erosion are some morphological operations that are able to mimic these historical transformations.

### 5.0.2 Dilation and Erosion

Dilation and erosion are two examples of morphological operations. Dilation adds pixels to the boundaries of objects in an image, while erosion removes pixels on object boundaries. The number of pixels added or removed from the objects in an image depends on the size and shape of the structuring element used to process the image. In the morphological dilation and erosion operations, the state of any given pixel in the output image is determined by applying a rule to the corresponding pixel and its neighbors in the input image. The rule used to process the pixels defines the operation as a dilation or an erosion. [3]



Figure 3: Flower Sample

Dilation - The value of the output pixel is the maximum value of all the pixels

in the input pixel's neighborhood. In a binary image, if any of the pixels is set to the value 1, the output pixel is set to 1. See figure 4.



Figure 4: Dilated flower

Erosion - The value of the output pixel is the minimum value of all the pixels in the input pixel's neighborhood. In a binary image, if any of the pixels is set to 0, the output pixel is set to 0. See figure 5.



Figure 5: Eroded flower

## 5.1 Adding Random Noise

Adding random noise to an image involves first generating an image of same size as the original image with pixel values from a random distribution. In our case,

we choose to use the normal distribution. Then we add the generated image to the original image (which one can think of as an overlay). See figure 6 below.

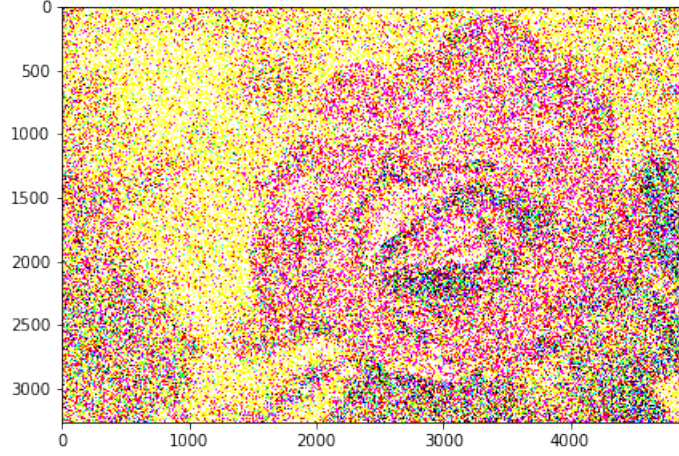


Figure 6: Noisy flower

## 5.2 Bubble Analysis

Given two matrix babbings of two images, we need a way to find the distance between them. This is accomplished by finding the distance between each bubble, and then combining the distance of all the bubbles. The distance between two bubbles can be found using a variety of statistical distances, summarized in the next section. Once we have a distance  $d$  that computes the distance between two bubbles, the total distance between the matrices is given as the  $L^1$  norm of the bubble distances. If  $P_{ij}$  are the bubbles of the first matrix and  $Q_{ij}$  are the bubbles of the second, then the total distance is given by

$$\sum_{i,j} d(P_{ij}, Q_{ij}).$$

It is possible that some bubbles would not have any data in them, which could be because the data is missing or otherwise does not exist in that area. In that case, there is no distribution for that bubble, so it is not included in the sum to calculate the total distance.

## 5.3 Distance Metrics

For comparing two bubbles, we considered several distance metrics.

### 5.3.1 Bhattacharyya Distance

The Bhattacharyya distance is a measure of distance between probability distributions related to the Hellinger distance. Like the Hellinger distance, it is a metric, satisfying the triangle inequality. The Bhattacharyya distance between two probability distributions  $P$  and  $Q$  with distribution functions  $p(x)$  and  $q(x)$  is given by [4]

$$B(P, Q) = -\log \int \sqrt{p(x)q(x)} dx.$$

This is similar to the Hellinger distance, which is

$$H(P, Q) = \left(1 - \int \sqrt{p(x)q(x)} dx\right)^{1/2}.$$

The relationship between the two of them is given by

$$B(P, Q) = -\log(1 - H^2(P, Q)).$$

For a multivariate normal distribution, we can use the value of  $H(P, Q)$  given above to derive

$$\begin{aligned} B(P, Q) = -\log & \frac{\det \Sigma_1^{1/4} \det \Sigma_2^{1/4}}{\det \left(\frac{\Sigma_1 + \Sigma_2}{2}\right)^{1/2}} \\ & + \frac{1}{8}(\mu_1 - \mu_2)^T \left(\frac{\Sigma_1 + \Sigma_2}{2}\right)^{-1} (\mu_1 - \mu_2). \end{aligned}$$

Note that if  $B$  and  $H$  are small, then to first order we have  $B(P, Q) \approx H^2(P, Q)$ .

### 5.3.2 Hellinger Distance

The Hellinger distance is another measure of the distance from one probability distribution to another. It is a metric, meaning that it is symmetric in its arguments and satisfies the triangle inequality. The Hellinger distance between two probability distributions is given by

$$H^2(P, Q) = 1 - \int \sqrt{f(x)g(x)} dx.$$

For multivariate normal distributions  $P$  and  $Q$ , where  $P$  has mean  $\mu_1$  and covariance matrix  $\Sigma_1$  and  $Q$  has mean  $\mu_2$  and covariance matrix  $\Sigma_2$ , the Hellinger distance is given by [4]

$$\begin{aligned} H^2(P, Q) = 1 - & \frac{\det \Sigma_1^{1/4} \det \Sigma_2^{1/4}}{\det \left(\frac{\Sigma_1 + \Sigma_2}{2}\right)^{1/2}} \\ & \exp -\frac{1}{8}(\mu_1 - \mu_2)^T \left(\frac{\Sigma_1 + \Sigma_2}{2}\right)^{-1} (\mu_1 - \mu_2) \end{aligned}$$



### 5.3.3 KL Divergence

The KL divergence is a measure of the distance from one probability distribution to another. Given a probability distribution  $P$  and another distribution  $Q$ , the KL divergence from  $P$  to  $Q$  can be thought of as a measure of how well the distribution  $Q$  approximates  $P$ . It is defined by

$$D_{KL}(P\|Q) = \int_{-\infty}^{\infty} p(x) \log \left( \frac{p(x)}{q(x)} \right) dx,$$

where  $p(x)$  and  $q(x)$  are the probability density functions of  $p$  and  $q$ . Since this distance is not symmetric in  $P$  and  $Q$ , it is not a true metric. Instead, we can use the *symmetrized KL divergence*, defined by

$$D_{KLs}(P, Q) = \frac{1}{2}(D_{KL}(P\|Q) + D_{KL}(Q\|P)).$$

Since we are approximating the data using normal distributions, we need a formula for the KL divergence from one normal distribution to another. Let  $P$  be a distribution with mean  $\mu_1$  and covariance matrix  $\Sigma_1$ , and  $Q$  a distribution with mean  $\mu_2$  and covariance matrix  $\Sigma_2$ . Then the probability density functions for  $P$  and  $Q$  are

$$p(x) = \frac{1}{\sqrt{(2\pi)^n \det \Sigma_1}} e^{-\frac{1}{2}(x-\mu_1)^T \Sigma_1^{-1} (x-\mu_1)}$$

$$q(x) = \frac{1}{\sqrt{(2\pi)^n \det \Sigma_2}} e^{-\frac{1}{2}(x-\mu_2)^T \Sigma_2^{-1} (x-\mu_2)}$$

and the  $\log(p(x)/q(x))$  is given by

$$\log \left( \frac{p(x)}{q(x)} \right) = \frac{1}{2} \log \frac{\det \Sigma_2}{\det \Sigma_1} + \frac{1}{2} (x - \mu_2)^T \Sigma_2^{-1} (x - \mu_2) - \frac{1}{2} (x - \mu_1)^T \Sigma_1^{-1} (x - \mu_1).$$

The quantity that we are looking for is the integral of  $p(x) \log(p(x)/q(x))$ . To make it easier to integrate, we can shift  $x$  by  $\mu_1$  and let  $\Delta\mu = \mu_2 - \mu_1$  to get

$$p(x) \log \frac{p(x)}{q(x)} = \frac{1}{2\sqrt{(2\pi)^n \det \Sigma_1}} \left( \log \frac{\det \Sigma_2}{\det \Sigma_1} + (x - \Delta\mu)^T \Sigma_2^{-1} (x - \Delta\mu) - x^T \Sigma_1^{-1} x \right) e^{-\frac{1}{2} x^T \Sigma_1^{-1} x}$$

Since any term with a single  $x$  is odd and goes to zero in the integral, the KL divergence is equal to

$$D_{KL}(P\|Q) = \frac{1}{2\sqrt{(2\pi)^n \det \Sigma_1}} \int (C + x^T (\Sigma_2^{-1} - \Sigma_1^{-1}) x) e^{-\frac{1}{2} x^T \Sigma_1^{-1} x} d^n x,$$

where

$$C = \log \frac{\det \Sigma_2}{\det \Sigma_1} + \Delta \mu^T \Sigma_2^{-1} \Delta \mu$$

To evaluate the integral, we need to make use of the Gaussian integral

$$\int e^{-\frac{1}{2}x^T A x} d^n x = \sqrt{\frac{(2\pi)^n}{\det A}},$$

for a symmetric matrix  $A$ , as well as the the following integral (see [5], pages 14-16):

$$\int x_i x_j e^{-\frac{1}{2}x^T A x} d^n x = (A^{-1})_{ij} \sqrt{\frac{(2\pi)^n}{\det A}}$$

Given another symmetric matrix  $B$ , we can take linear combinations of the above integral to derive

$$\int x^T B x e^{-\frac{1}{2}x^T A x} d^n x = \sqrt{\frac{(2\pi)^n}{\det A}} \text{Tr } A^{-1} B$$

Using these two integrals, we can calculate the KL divergence to be

$$\begin{aligned} D_{KL}(P\|Q) &= \frac{1}{2}C + \frac{1}{2} \text{Tr}(\Sigma_1(\Sigma_2^{-1} - \Sigma_1^{-1})) = \frac{1}{2}C - \frac{n}{2} + \frac{1}{2} \text{Tr}(\Sigma_1 \Sigma_2^{-1}) \\ &= \frac{1}{2} \left( \log \frac{\det \Sigma_2}{\det \Sigma_1} + \Delta \mu^T \Sigma_2^{-1} \Delta \mu + \text{Tr}(\Sigma_1 \Sigma_2^{-1}) - n \right) \end{aligned}$$

### 5.3.4 Fisher Information

One way that the KL divergence can be made symmetric is by symmetrizing it as described above, to get the symmetric KL divergence. Another way to get a metric from the KL divergence is to look at the Hessian with respect to  $Q$  for a fixed  $P$ . Doing so will give a matrix which can be interpreted as a Riemannian metric tensor. This metric tensor can then be used to get a metric on the manifold of probability distributions, and the metric is called the Fisher information metric. If the probability distributions depend on some variables  $\theta_i$ , then the KL divergence is

$$D_{KL}(P(\theta^0)\|P(\theta)) = - \int p(x; \theta^0) \log \frac{p(x; \theta)}{p(x; \theta^0)} dx$$

Taking the derivative of the integrand with respect to first  $\theta_i$ , we get

$$-p(\theta^0) \frac{1}{p(\theta)} \frac{dp}{d\theta_i}.$$

(At this point, if we let  $\theta = \theta_0$ , when we integrate over all  $x$  we get 0, as expected.) Then taking a second derivative with respect to  $\theta_j$  we get

$$p(\theta^0) \frac{1}{p(\theta)^2} \frac{dp}{d\theta_i} \frac{dp}{d\theta_j} - p(\theta^0) \frac{1}{p(\theta)} \frac{d^2 p}{d\theta_i d\theta_j}$$

When we let  $\theta = \theta_0$  and integrate over all  $x$ , the second term goes to zero, and we are left with

$$\int \frac{1}{p(x; \theta)} \frac{dp}{d\theta_i} \frac{dp}{d\theta_j} dx.$$

Using the fact that

$$\frac{d \log p(x; \theta)}{d\theta_i} = \frac{1}{p(x; \theta)} \frac{dp(x; \theta)}{d\theta_i},$$

we can recover the more familiar form

$$\int \frac{d \log p(x; \theta)}{d\theta_i} \frac{d \log p(x; \theta)}{d\theta_j} p(x; \theta) dx.$$

On the space of single-variate Gaussian distributions, there are only two parameters, the mean  $\mu$  and the variance  $\sigma$ . That means that the metric tensor is a  $2 \times 2$  matrix, and it happens to equal

$$\frac{1}{\sigma^2} \begin{pmatrix} 1 & 0 \\ 0 & 2 \end{pmatrix}.$$

This is the same as the metric for the Poincare model of the hyperbolic plane, so under the Fisher information metric, distances between two normal distributions are the same as hyperbolic distances in the plane.

One potential future application of matrix bubbling that we have not yet considered is to time series data. If the data is a continuous function of time, then we would be interested in looking at the derivative. The Fisher information metric can be of use in calculating this derivative. Given  $P(\theta)$  where the  $\theta$  parameters are functions of the time variable  $t$ , the derivative would be given by

$$\frac{d(P(\theta(t + \Delta t)), P(\theta(t)))}{\Delta t}$$

Letting  $\Delta\theta = \theta(t + \Delta t) - \theta(t)$ , this is equal to

$$\frac{d(P(\theta + \Delta\theta), P(\theta))}{\Delta t}.$$

The distance here could be any distance, but since the change in  $\theta$  would be small for a single time step, the Fisher information metric could be used as a good approximation. If the metric tensor is  $g_{ij}(\theta)$ , then the squared distance between  $P(\theta + \Delta\theta)$  and  $P(\theta)$  is

$$g_{ij}(\theta) \Delta\theta_i \Delta\theta_j.$$

So the derivative is then equal to

$$\frac{\sqrt{g_{ij}(\theta) \Delta\theta_i \Delta\theta_j}}{\Delta t}$$

## 6 Results

Below is a grid of images with morphological operators applied to them.

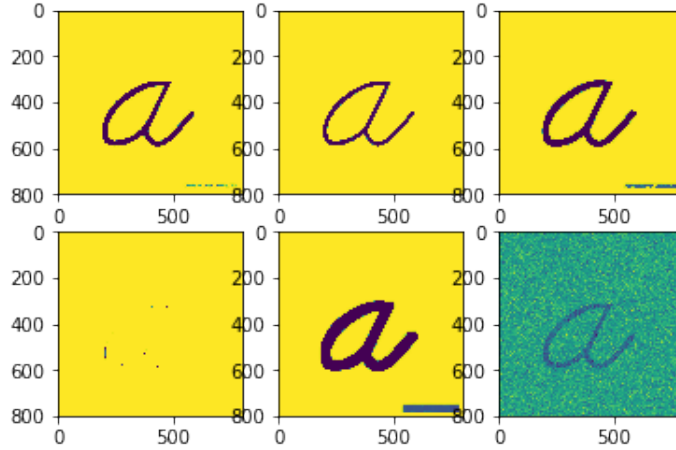


Figure 7: Morphological Operators Applied on an Image

The first image is the original image. Sweeping from left to right then downwards, we see that the second image (A) is an erosion, the third (B) is a dilation, the fourth (C) is a stronger erosion, the fifth (D) is a stronger dilation, and the sixth (E) has added random noise.

We compared each of the five morphed images with the original using both the Hellinger distance and the KL divergence. The distances are recording in the following table. Some of the comparisons did not work due to the code crashing so we have put a '-' in the space to denote that we should ignore it for now.

	A	B	C	D	E
Hellinger	-8.24	24.20	0.25	51.97	-
KL	3092.31	-	3101.11	-	-

Figure 8: Comparison of Morphed Letters with Original Letter

We also compared each of the five morphed images with the letter 'B'. We did this because we expect the distances to be greater since they are not the same letter. The distances are recorded in the following table.

For the Hellinger distance, we can see that the distances in figure 8 are generally larger than those in figure 9. This makes sense since similar images should give a smaller distance. For the same reason, the highly eroded image in general had a larger distance compared to the other morphed images. We also see that the



ALPHABETPRINTABLES.ORG

Figure 9: Cursive Letter ‘B’

	A	B	C	D	E
Hellinger	14.72	45.15	1.26	54.18	-
KL	35.55	-	330.40	-	-

Figure 10: Comparison of Morphed Letters with Different Letter

scores for the strongly eroded images are lower than the others since there are not much pixels to compare since we throw out grids with no pixels. Also, we note that the first comparison gives us a negative number, which should not be possible.

For the KL divergence, there seem to be much more errors and the trend seems to be the opposite of what we expect. We see that similar images have high divergence while different images have low divergence. This means that we should probably not use the KL divergence.

## 7 Future Work

Our progress so far has been to understand and come up with an initial implementation for the matrix bubbling technique. The next steps are to investigate more applications of matrix bubbling, and specifically to see how well it performs. We would also like to improve our code so it can run faster and produce fewer errors.

## References

- [1] URL: <https://tex.stackexchange.com/questions/375285/draw-a-grid-of-bivariate-normal-distributions-in-tikz>.
- [2] URL: [https://docs.opencv.org/3.0-beta/doc/py\\_tutorials/py\\_imgproc/py\\_morphological\\_ops/py\\_morphological\\_ops.html](https://docs.opencv.org/3.0-beta/doc/py_tutorials/py_imgproc/py_morphological_ops/py_morphological_ops.html).
- [3] URL: <https://www.mathworks.com/help/images/morphological-dilation-and-erosion.html>.
- [4] Leandro Pardo. *Statistical Inference Based on Divergence Measures*.
- [5] Anthony Zee. *Quantum Field Theory in a Nutshell*.